

CIS 452 - Operating Systems Concepts

Nathan Bowman

Images taken from Silberschatz book

More Paging

More Paging

Physical memory and logical memory are split into frames and pages

Any page can be mapped to any frame

Mapping is transparent to process

To translate addresses, MMU and OS must work together (as has been the case to some degree for all of our schemes)

MMU handles actual translation, but OS is responsible for managing page table

There is a page table associated with each process

OS is responsible for filling page tables

OS must know about available frames to do its job

Keeps frame information in **frame table**

free-frame list


14

13

18

20

15



13

14

15

16

17

18

19


20

21

(a)

free-frame list

15



13

14

15

16

17

18

19

20

21

new-process page table

0	14
1	13
2	18
3	20

(b)

page 1


page 0

page 2

page 3

Memory access is protected by design

Physical address determined by page table, which is
controlled by OS



Paging can also make implementing shared memory simple

If OS puts same frame into page table of two (or more) processes, then they will both be allowed to access same memory

Note that it is a shared *frame* (physical memory), not a shared *page* (logical memory) that matters here

Shared pages can also be useful outside of IPC

Imagine 40 users all want to use same editor, and editor requires 150 KB of code and 50 KB of data in memory -
- total memory demand is 8,000 KB

Since code is read-only, it can be shared and only one copy need be kept in memory

ed 1
ed 2
ed 3
data 1

process P_1

3
4
6
1

page table
for P_1

ed 1
ed 2
ed 3
data 2

process P_2

3
4
6
7

page table
for P_2

ed 1
ed 2
ed 3
data 3

process P_3

3
4
6
2

page table
for P_3

0	
1	data 1
2	data 3
3	ed 1
4	ed 2
5	
6	ed 3
7	data 2
8	
9	
10	
11	

Total memory for previous example is now $40 \times 50 \text{ KB} + 150 \text{ KB} = 2,150 \text{ KB}$

Contrasted with previous 8,000 KB, this is a large savings

Benefits:

- No external fragmentation
- Maps well to backing store
- Simple protection
- Easy sharing of memory

Backing store will have similar issues as memory

Backing now store split into blocks of the same size as frames/pages, making management of backing store itself simpler

One downside is internal fragmentation

On average, each process will have $(\text{page_size}/2)$ bytes
of internal fragmentation

In terms of fragmentation, smaller pages are better

However, additional pages require additional storage overhead

Also, disks are generally more efficient when transferring large amounts of data

Page sizes in examples were obviously unreasonable

Typical memory for personal computers is several gigabytes, and page sizes are several kilobytes to megabytes (or sometimes gigabytes)

```
$ getconf PAGESIZE  
4096
```

Most OSes use paging