

CIS 452 - Operating Systems Concepts

Nathan Bowman

Images taken from Silberschatz book

Structure of Page Tables

Page tables are stored in memory because we do not have enough registers to store an entire table

Even in memory, storing large page tables can be prohibitively expensive

Most modern computers are 32- or 64-bit address architectures

The logical address space is then (2^{32}) to (2^{64}) bits

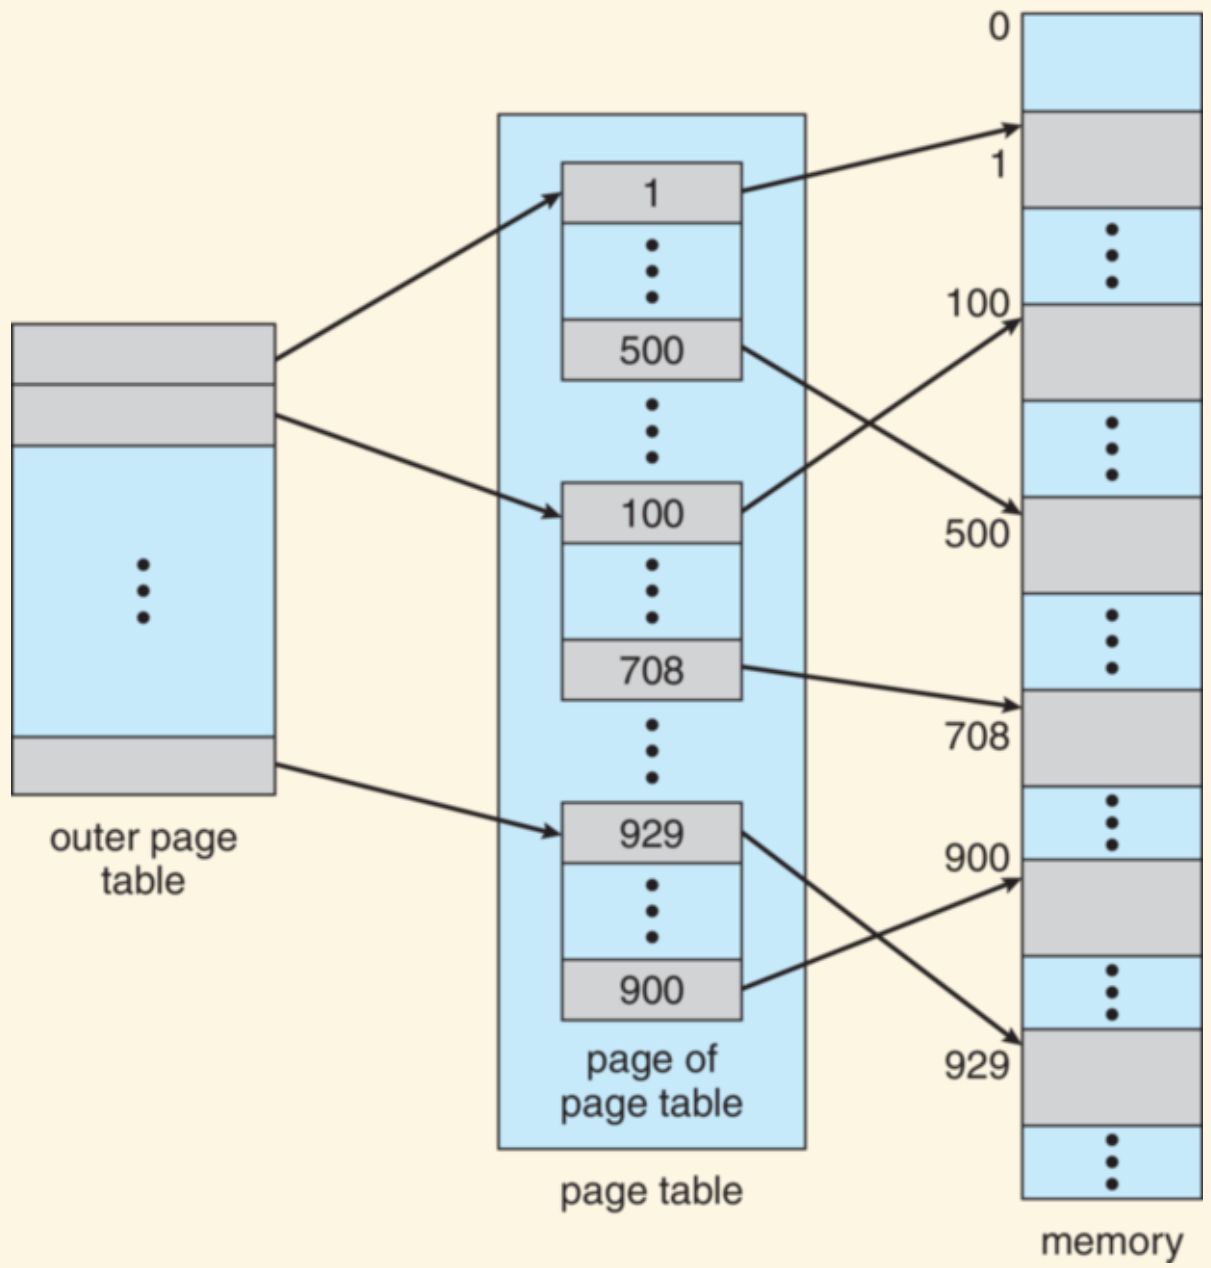
In a 32-bit system with page size of 4 KB (2^{12}), page table needs $(2^{32}/2^{12} = 2^{20})$ entries, which is more than 1 million

Assuming each entry is 4 bytes, the page table *for each process* could require up to 4 MB of physical memory

Note that, until now, the page table had been assumed to be stored contiguously in memory

Requiring 4 MB of physical memory is bad enough, but to keep the page tables contiguously in memory would be extremely inconvenient

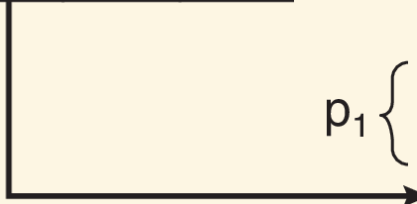
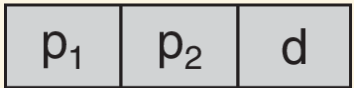
We can instead page the page table



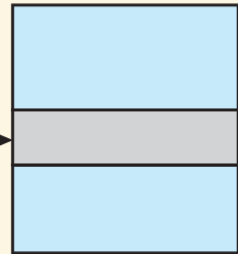
"Page number" portion of address is then split into two parts

First part now represents which "page of page table" is requested

logical address



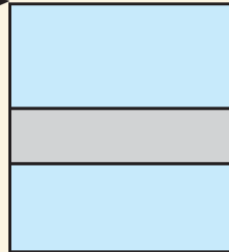
p_1 {



outer page
table



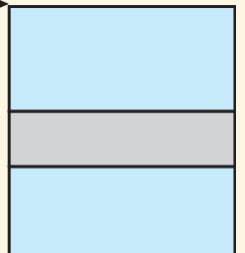
p_2 {



page of
page table

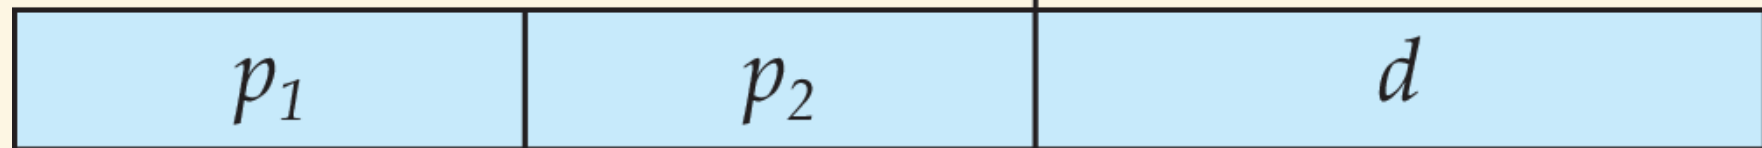


d {



page number

page offset



p_1

p_2

d

10

10

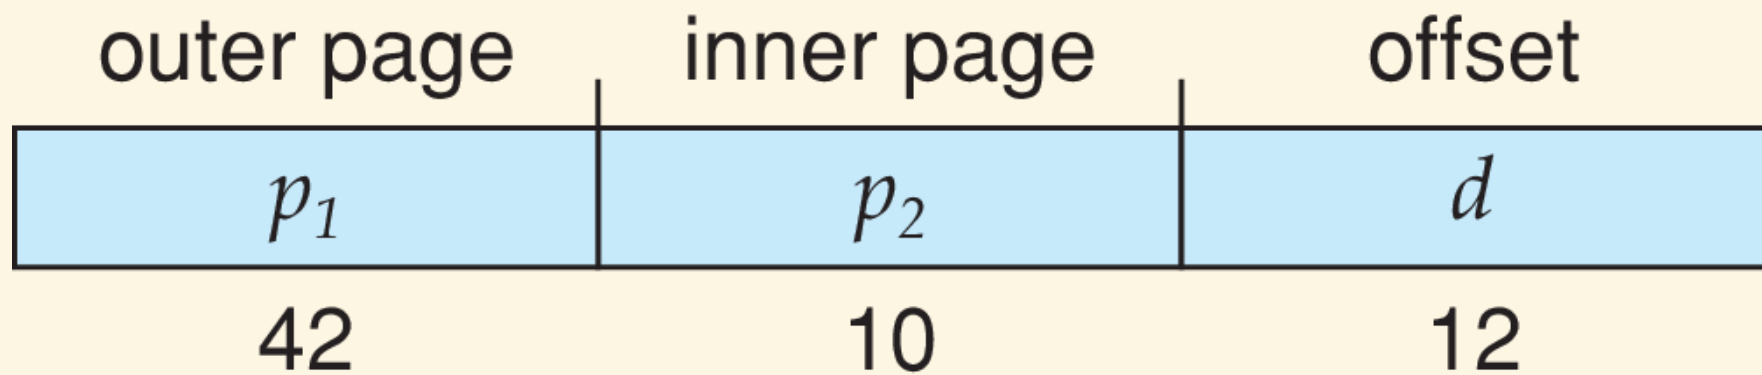
12

This is an example of **hierarchical paging** known as a **forward-mapped page table**

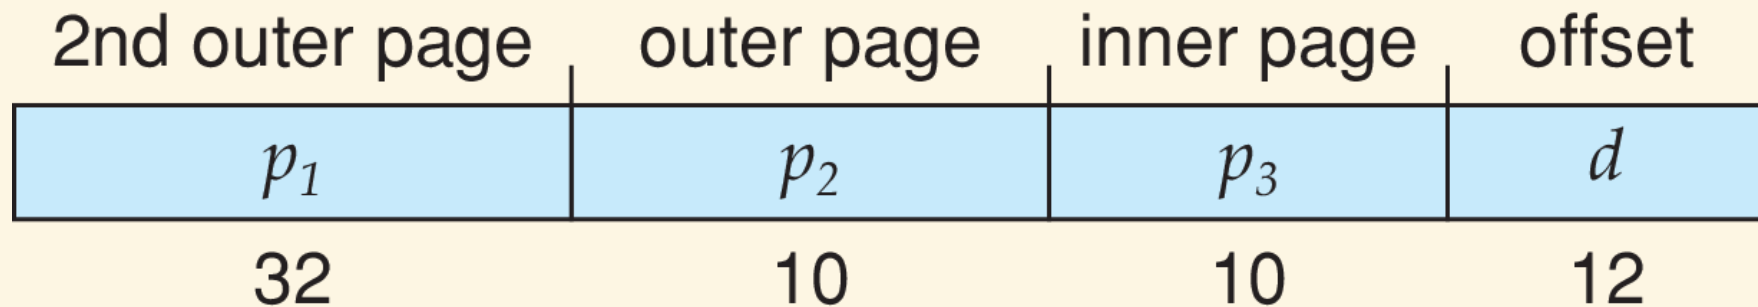
We need only store page table entries for pages that are mapped to physical frames

Large portion of virtual address space may go unused, so this could be substantial memory savings

For a 64-bit machine, this will not break the page table into small enough chunks to be manageable



Outer page is much too large
What if we repeat the same idea?



Even so, outer page table is still 16 GB

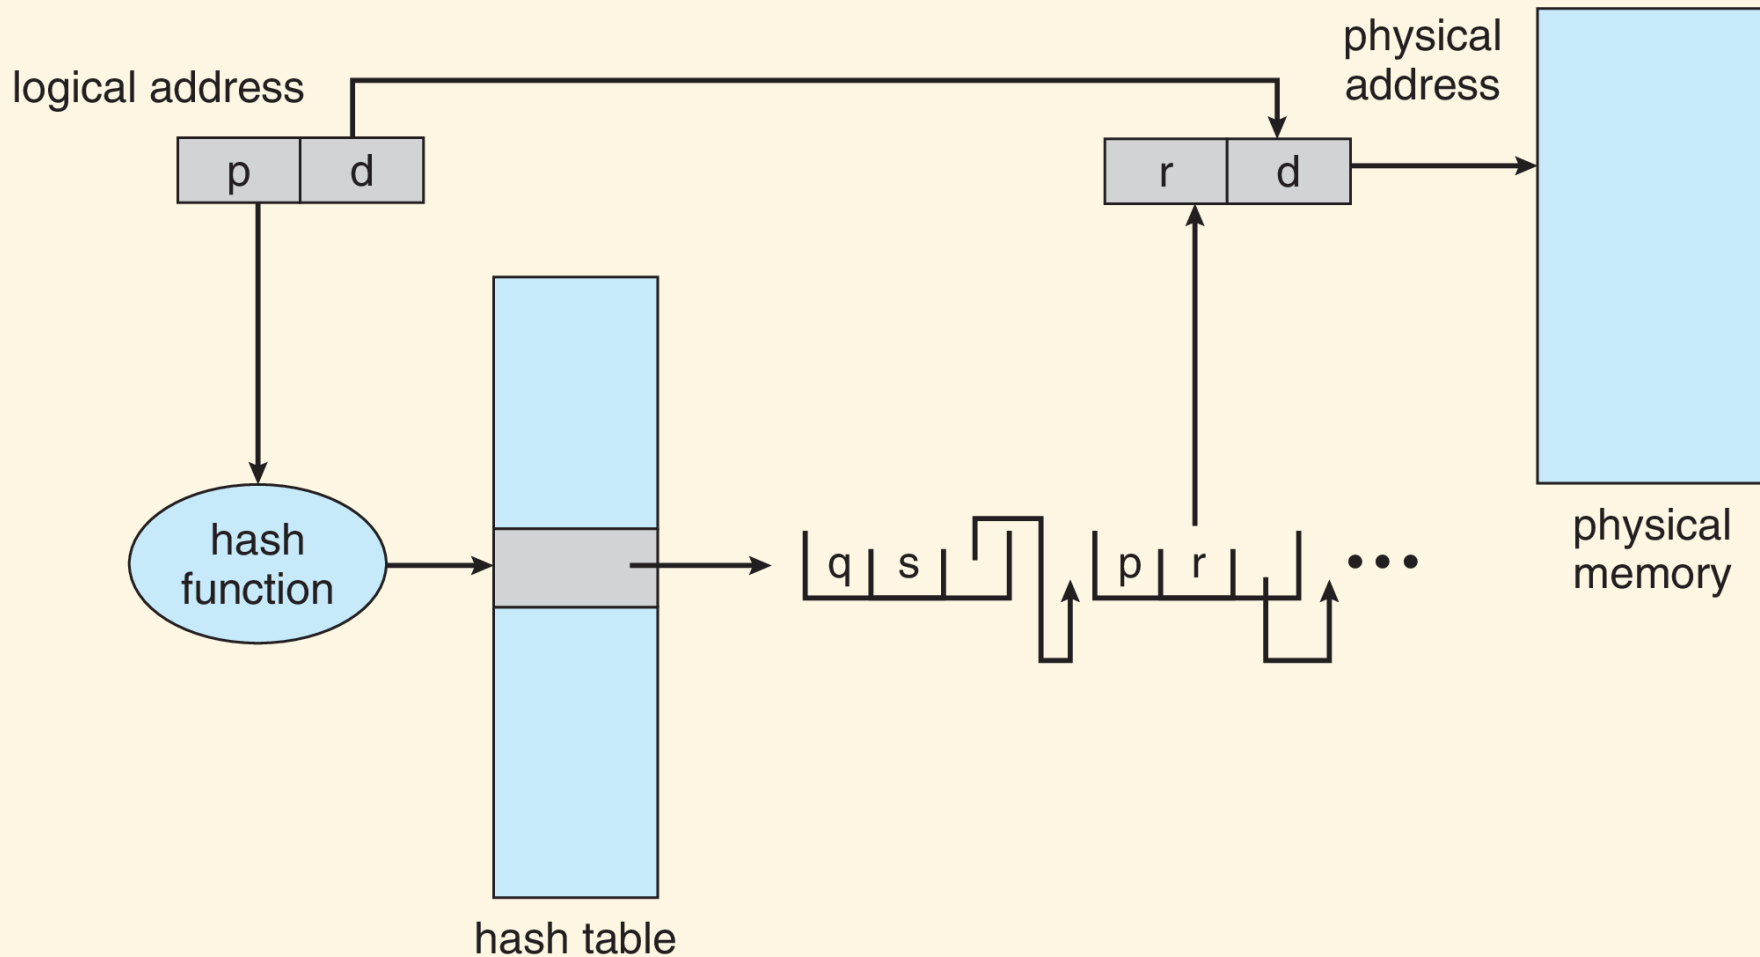
Could keep going, but the number of required memory accesses goes up each time

For larger address spaces, hashed page table is a better choice

Hashed page table is hash table whose key is page number

Values are linked lists of elements, where element contains

- page number
- corresponding frame number
- pointer to next element



In this setup, only need to store page table entries for pages that are mapped to physical frames

Large portion of virtual address space may go unused, so this could be substantial memory savings

One downside to hashed page tables is that each process still has its own page table

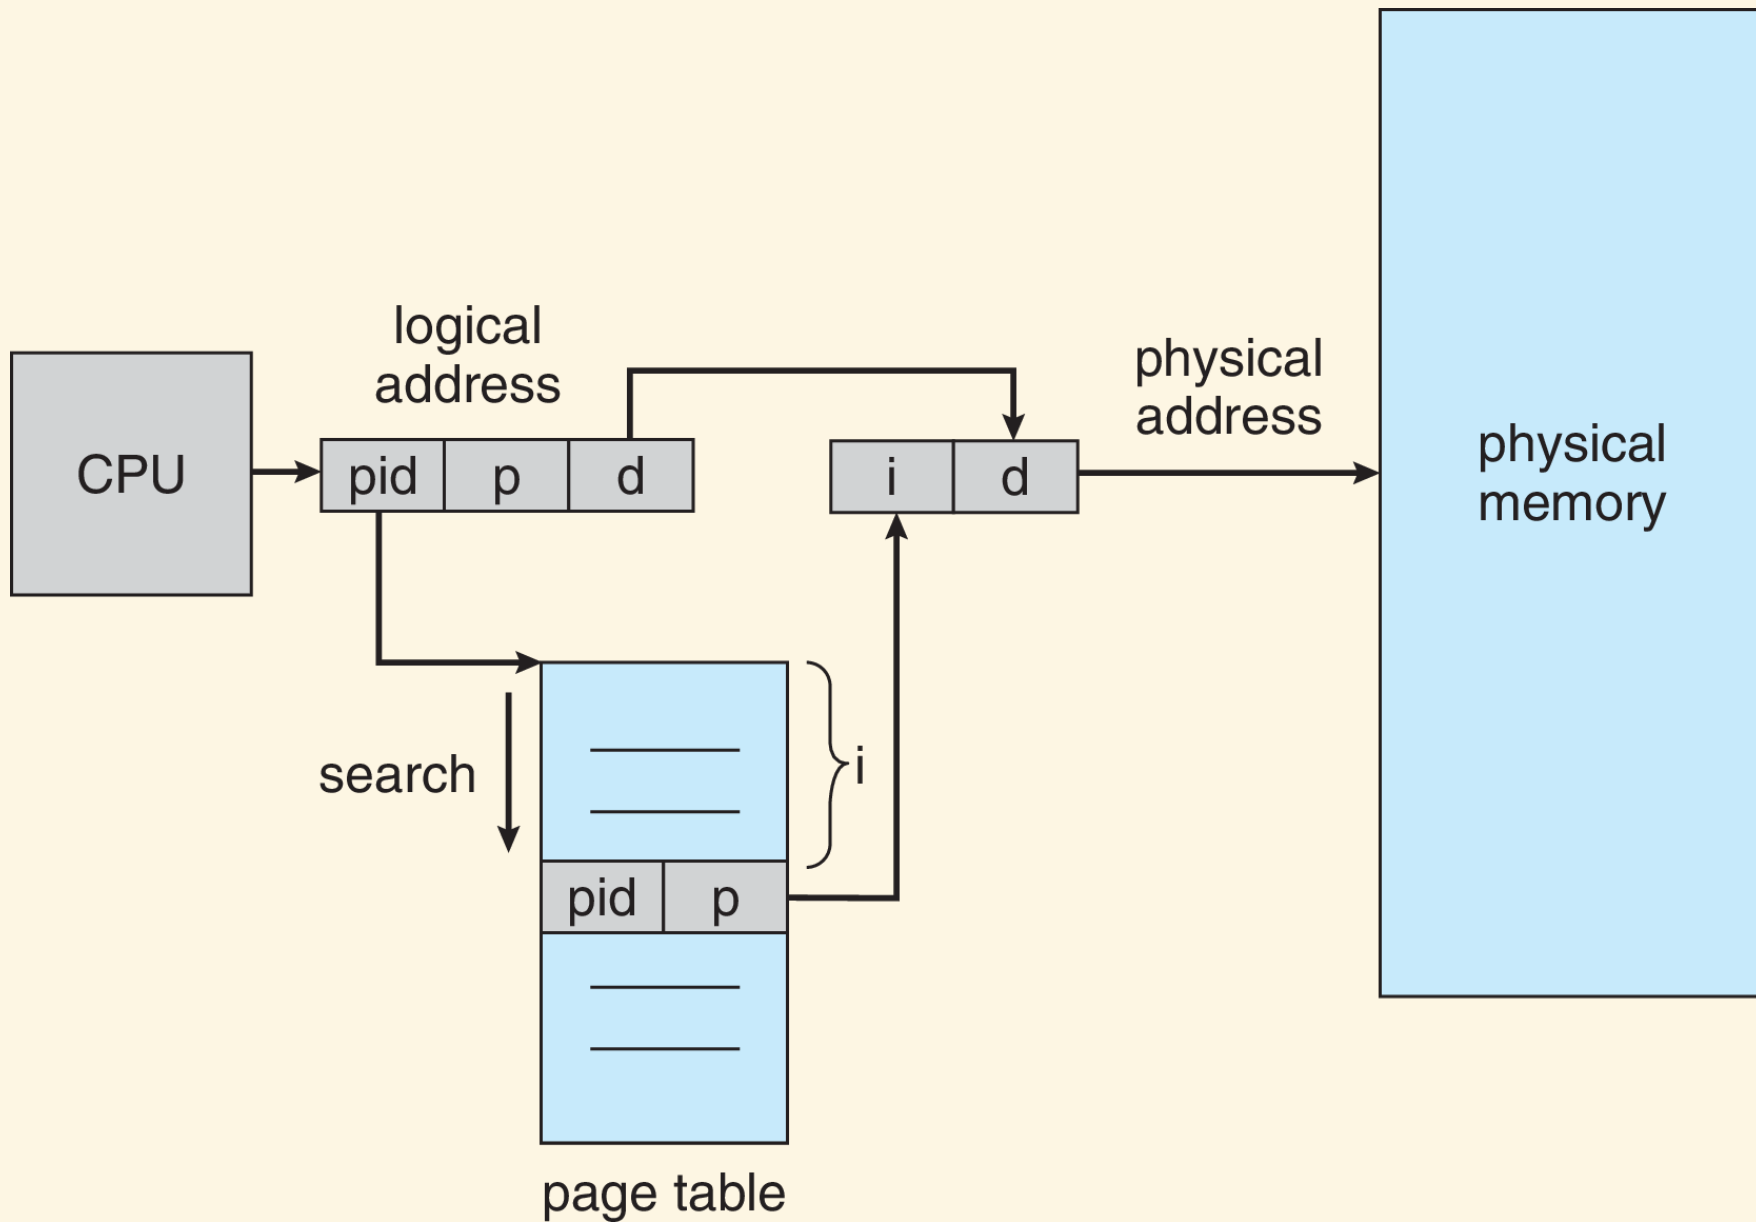
There may be many processes, but there are only so many frames of memory

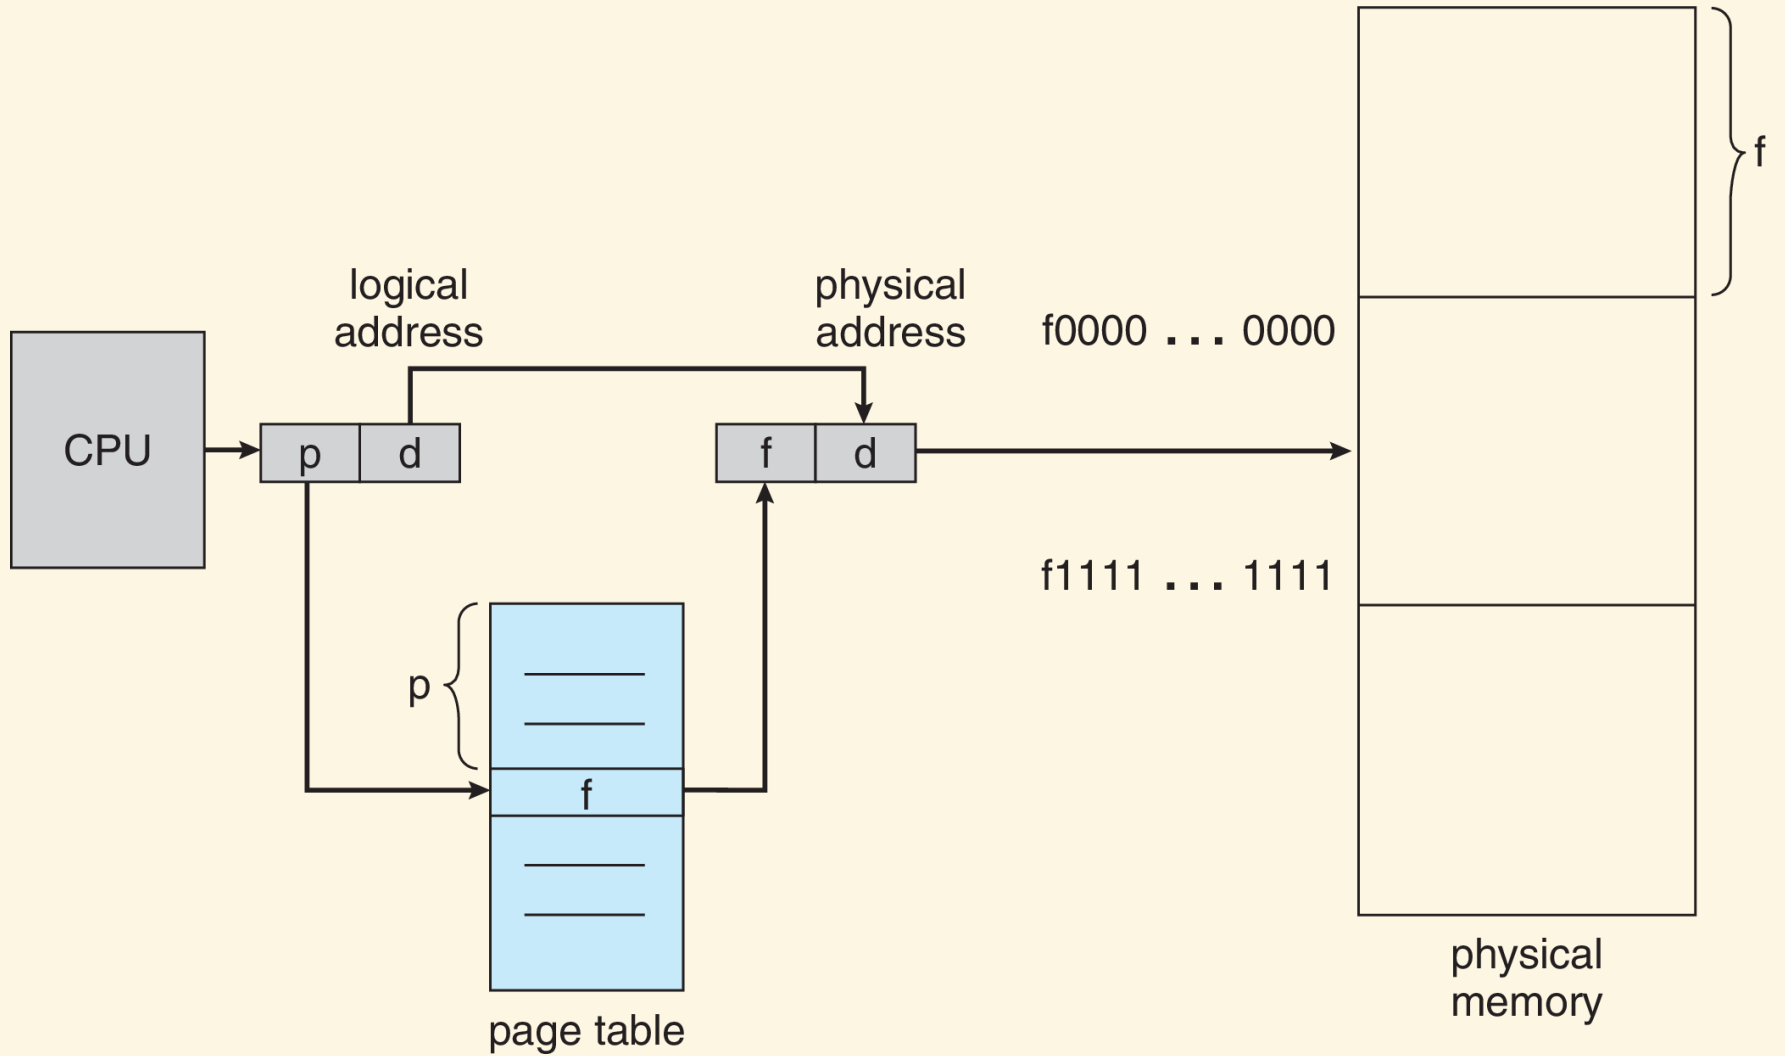
Inverted page table stores just one entry per frame

Information is kept about corresponding virtual address
for each frame

Only need one inverted page table system-wide, as
opposed to one per process

Because of this, entries of table must also track which
process owns each frame





This is great for saving memory, but it does not match
how addresses are generated

CPU generates logical addresses which must then be
mapped to physical addresses

Storing table in order of logical addresses makes lookup
easy, but storing by physical address gives us no clue
where to look

Searching every table entry for match would not be
feasible

Instead, use hash table again

Map logical address to linked list and search list until a match for **both** PID and page number is found

Difference from previous hash setup is that we now store each frame address just once, so there is one table per system, which is why storing PID is also necessary

Small memory footprint and much faster than searching entire table, though still requires more than one memory access

Note that regardless of which scheme is used, a TLB should still be used for efficiency

There are four basic ways to store page tables in memory

- direct storage
- hierarchical page tables
- hashed page tables
- inverted page table

These have tradeoffs between speed and memory footprint