

CIS 452 - Operating Systems Concepts

Nathan Bowman

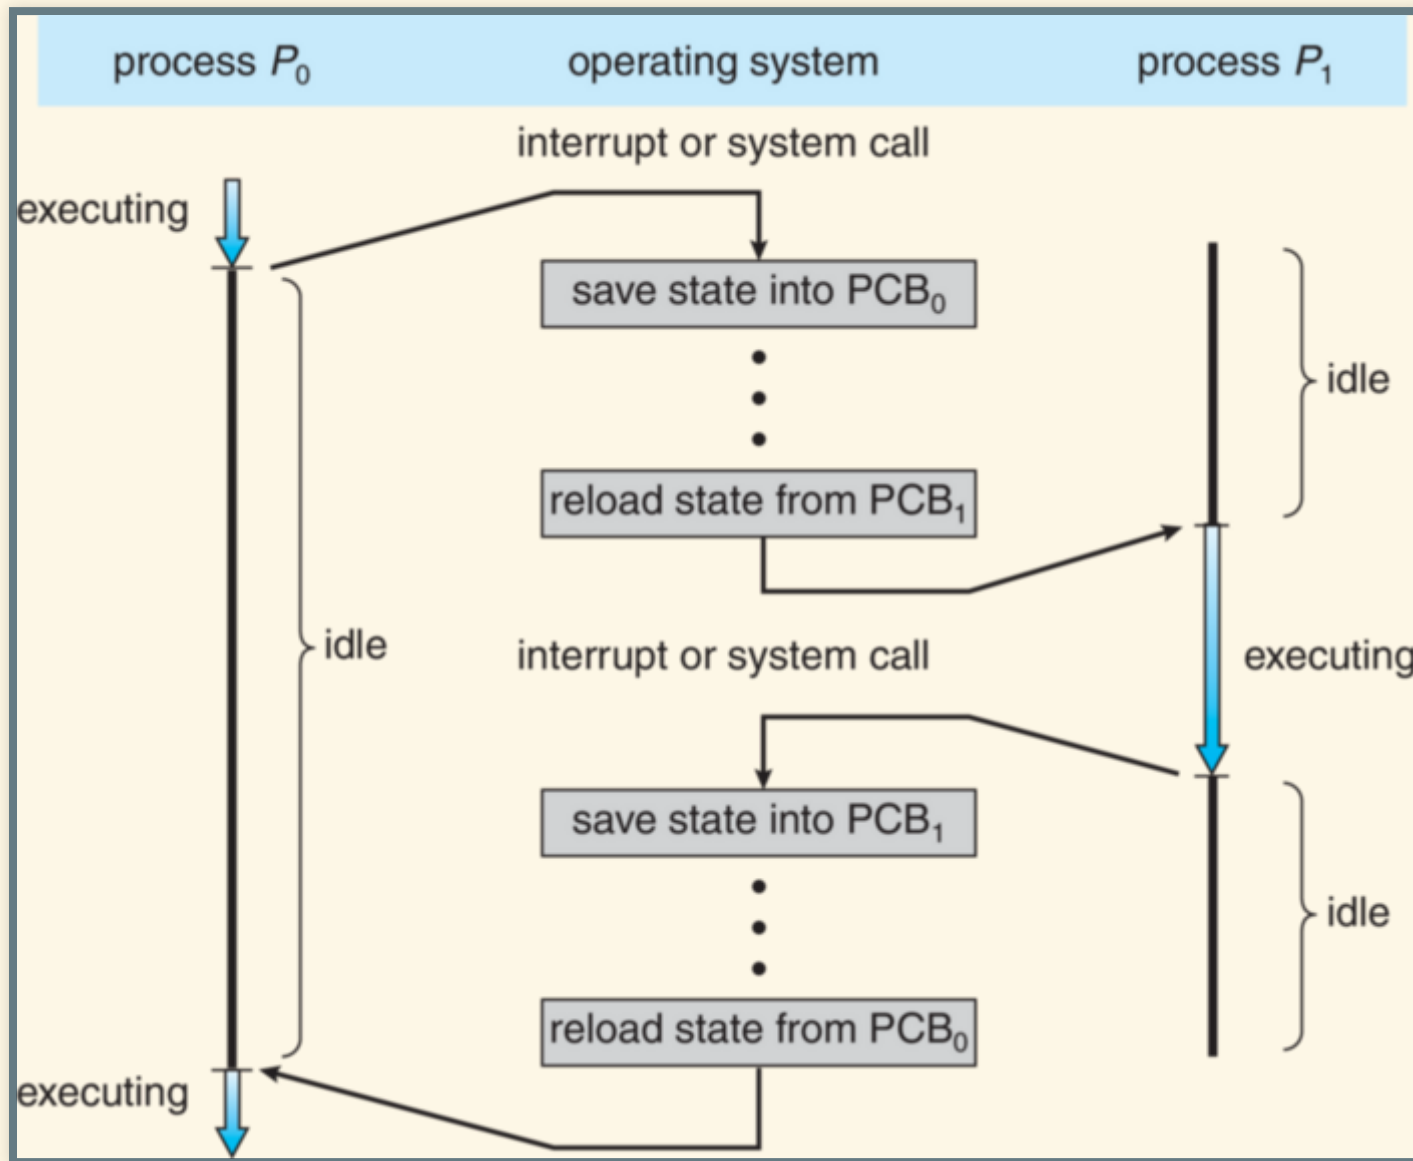
Images taken from Silberschatz book

---

Process Scheduling

OS regularly changes which process controls CPU

When running process changes, entire state must  
change with it



Another word for state is "context" -- this process is a  
**context switch**

Process's entire context (everything it needs to run correctly) is stored in PCB

Context switch is when OS saves state of running process into PCB, restores state of next process from its PCB, and puts that process on the CPU

Context switch metaphor is quite useful for real life

Think about moving from reading your textbook to responding to an email, for example

Time spent in a context switch is completely wasted

Duration of context switch depends on several factors:

- complexity of stored state
- hardware speed
- special instructions from architecture

Few milliseconds is a good approximation

OS be careful how it decides which processes to schedule

Must consider resource utilization (via multiprogramming) and time sharing (how reactive is computer?)

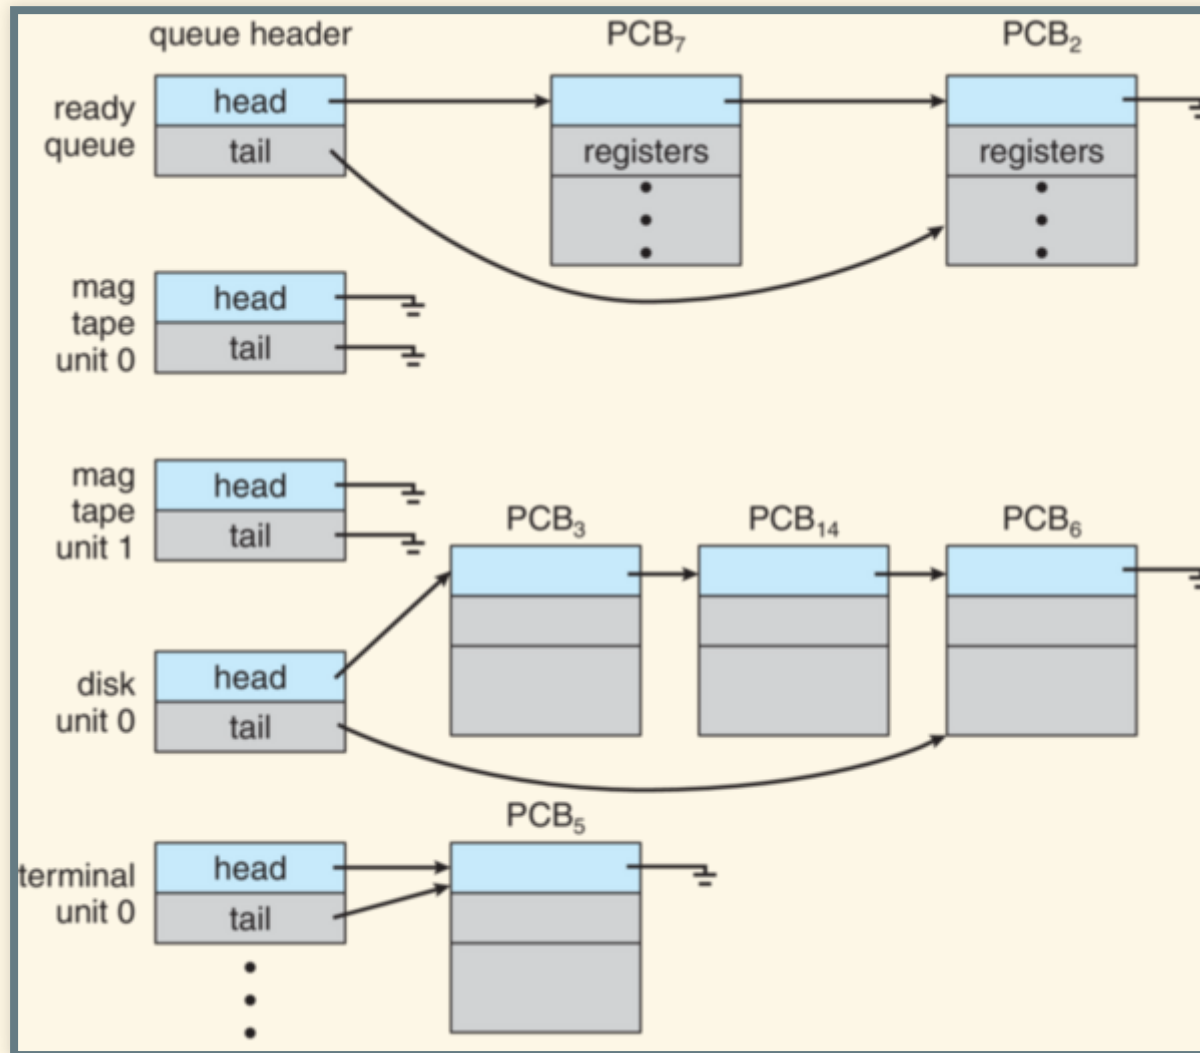
Correct balance depends on use case

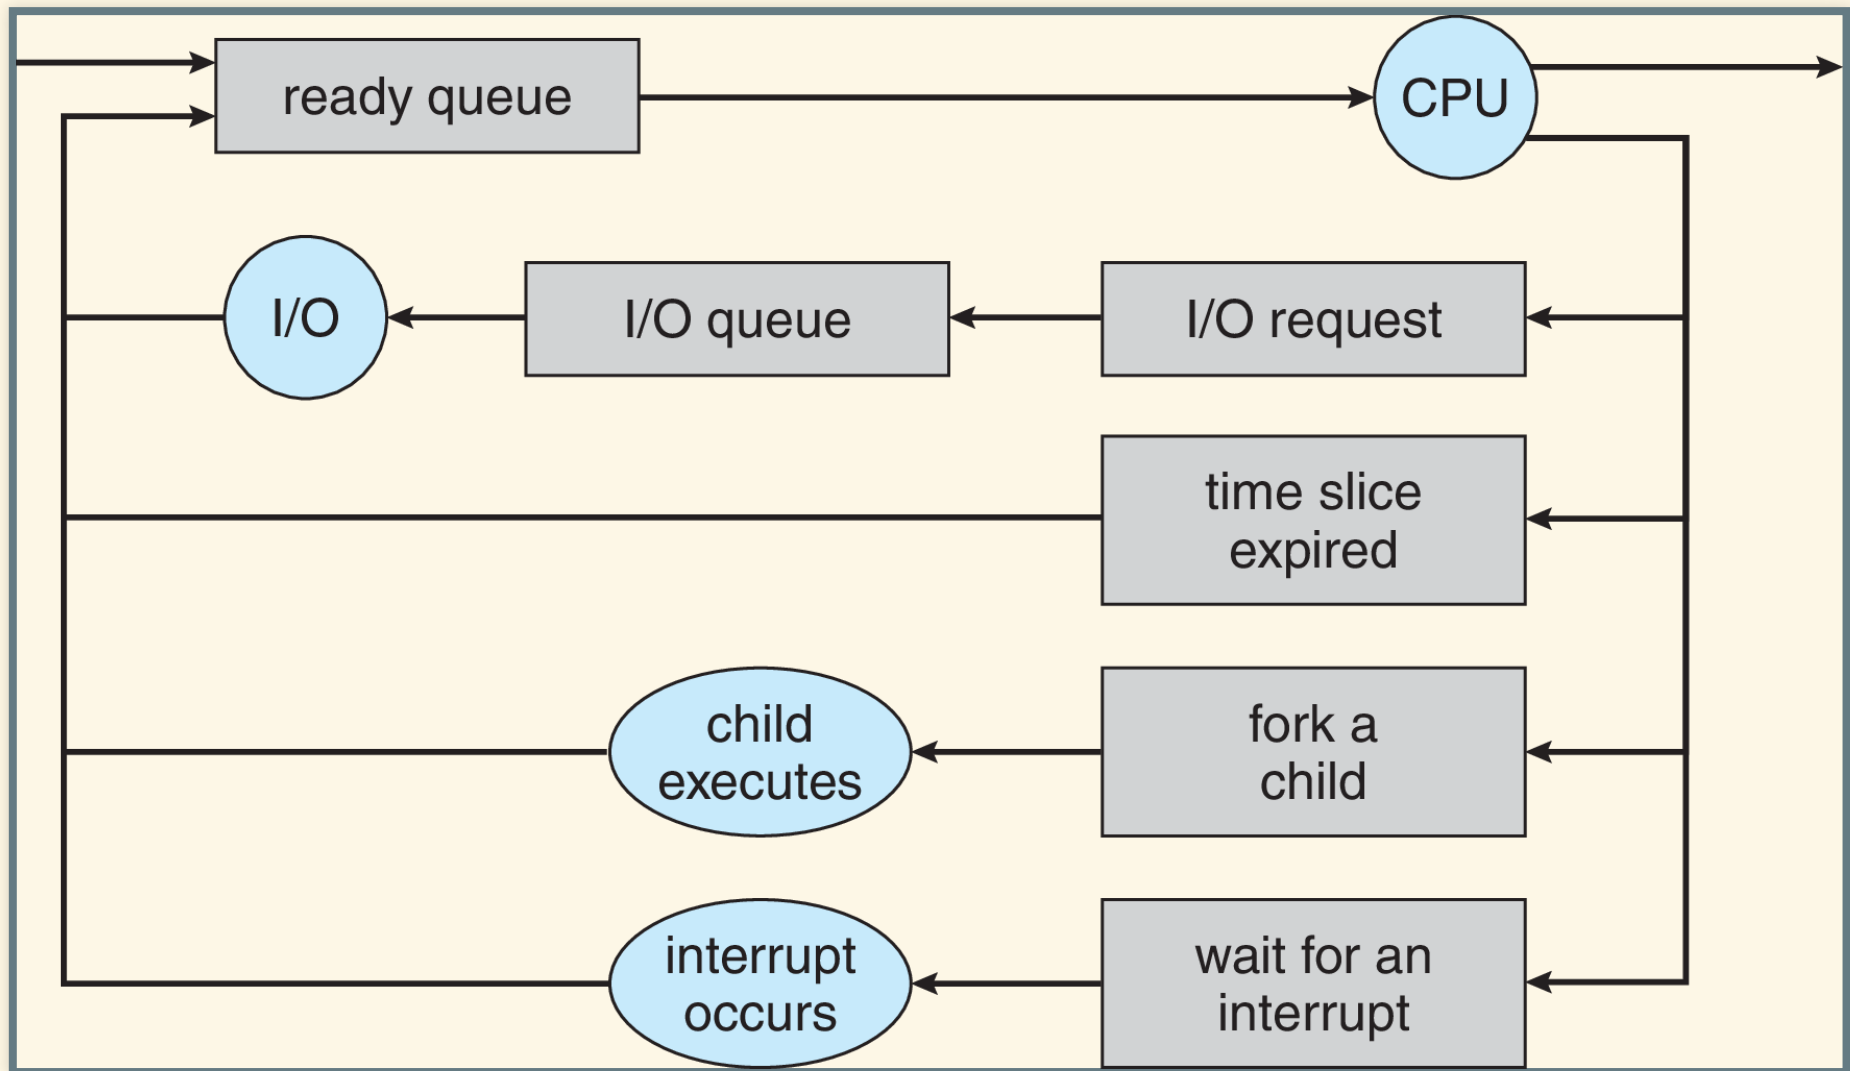
Process scheduler -- deciding who gets the CPU

Relies on several queues to keep track of process state

- job queue -- process waiting to get into main memory
- ready queue -- processes in main memory ready and waiting for CPU
- device queues -- processes waiting for I/O on a particular device (one queue per device)







## Levels of schedulers

- CPU (short-term) scheduler
- medium-term scheduler
- job (long-term) scheduler

CPU scheduler determines which process to run next

Scheduler is running every few milliseconds to keep  
system responsive

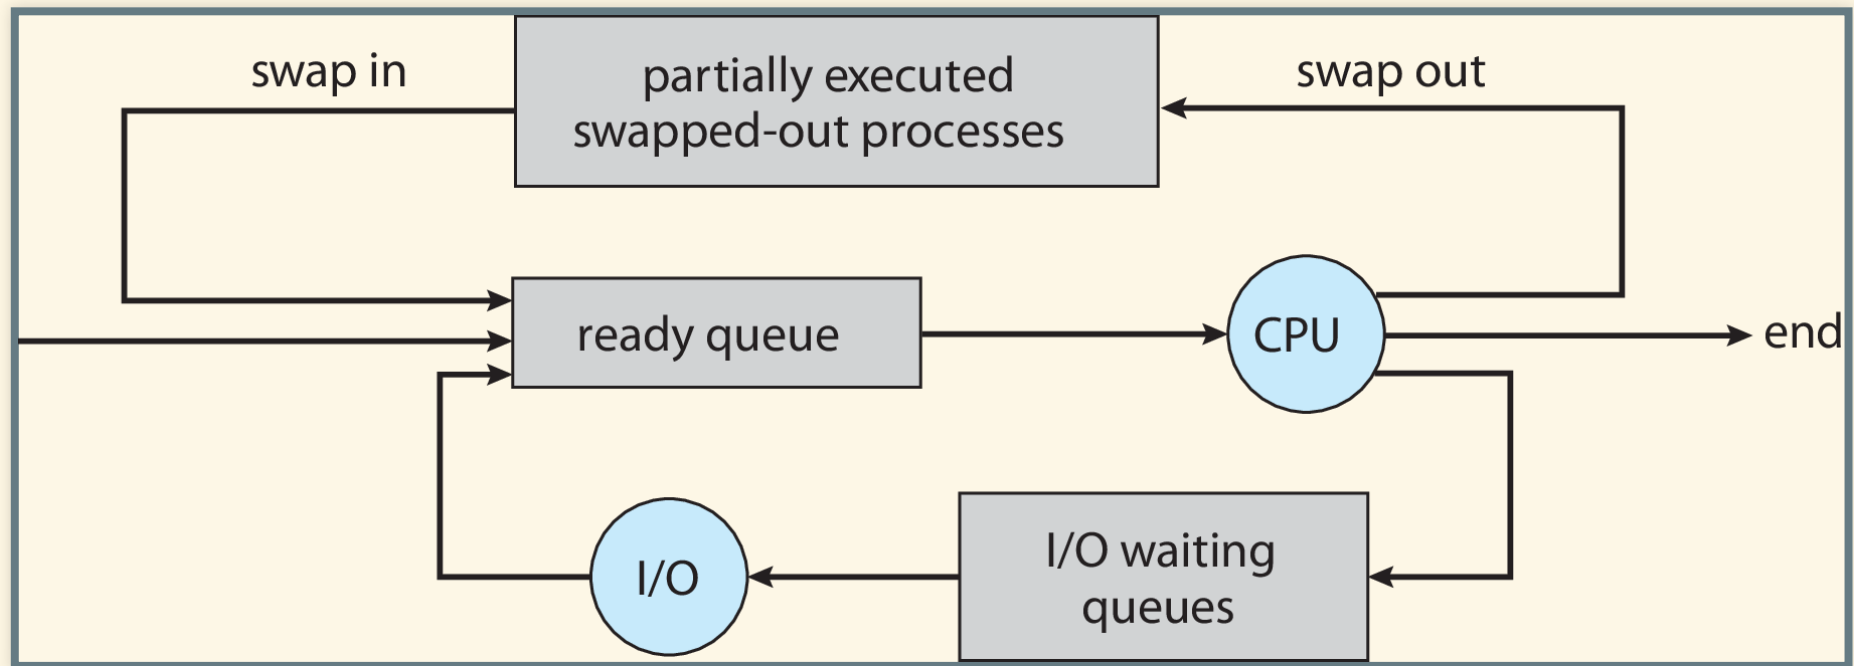
Must be **fast** (remember: time spent context switching  
is time wasted)

Medium-term scheduler is responsible for temporarily removing processes from memory as needed

Moving processes into and out of memory is **swapping**

Sometimes done to achieve better balance of CPU-bound and I/O-bound processes

Also done if memory is overtaxed (this is when you will notice your computer getting *very* slow)



Swapping is another nice metaphor for life

Memory is computer's place to do useful work

Running out of memory is like trying to cook in a tiny kitchen or work on an overcrowded desk

Job scheduler determines which processes get into memory

This applies more to batch systems, whereas medium-term scheduling applies more to time-sharing systems

Usually jobs in == jobs out, so this scheduler runs only every few minutes and can afford to be much slower

Tries to strike good balance between I/O-bound and CPU-bound processes