


CIS 452 - Operating Systems Concepts

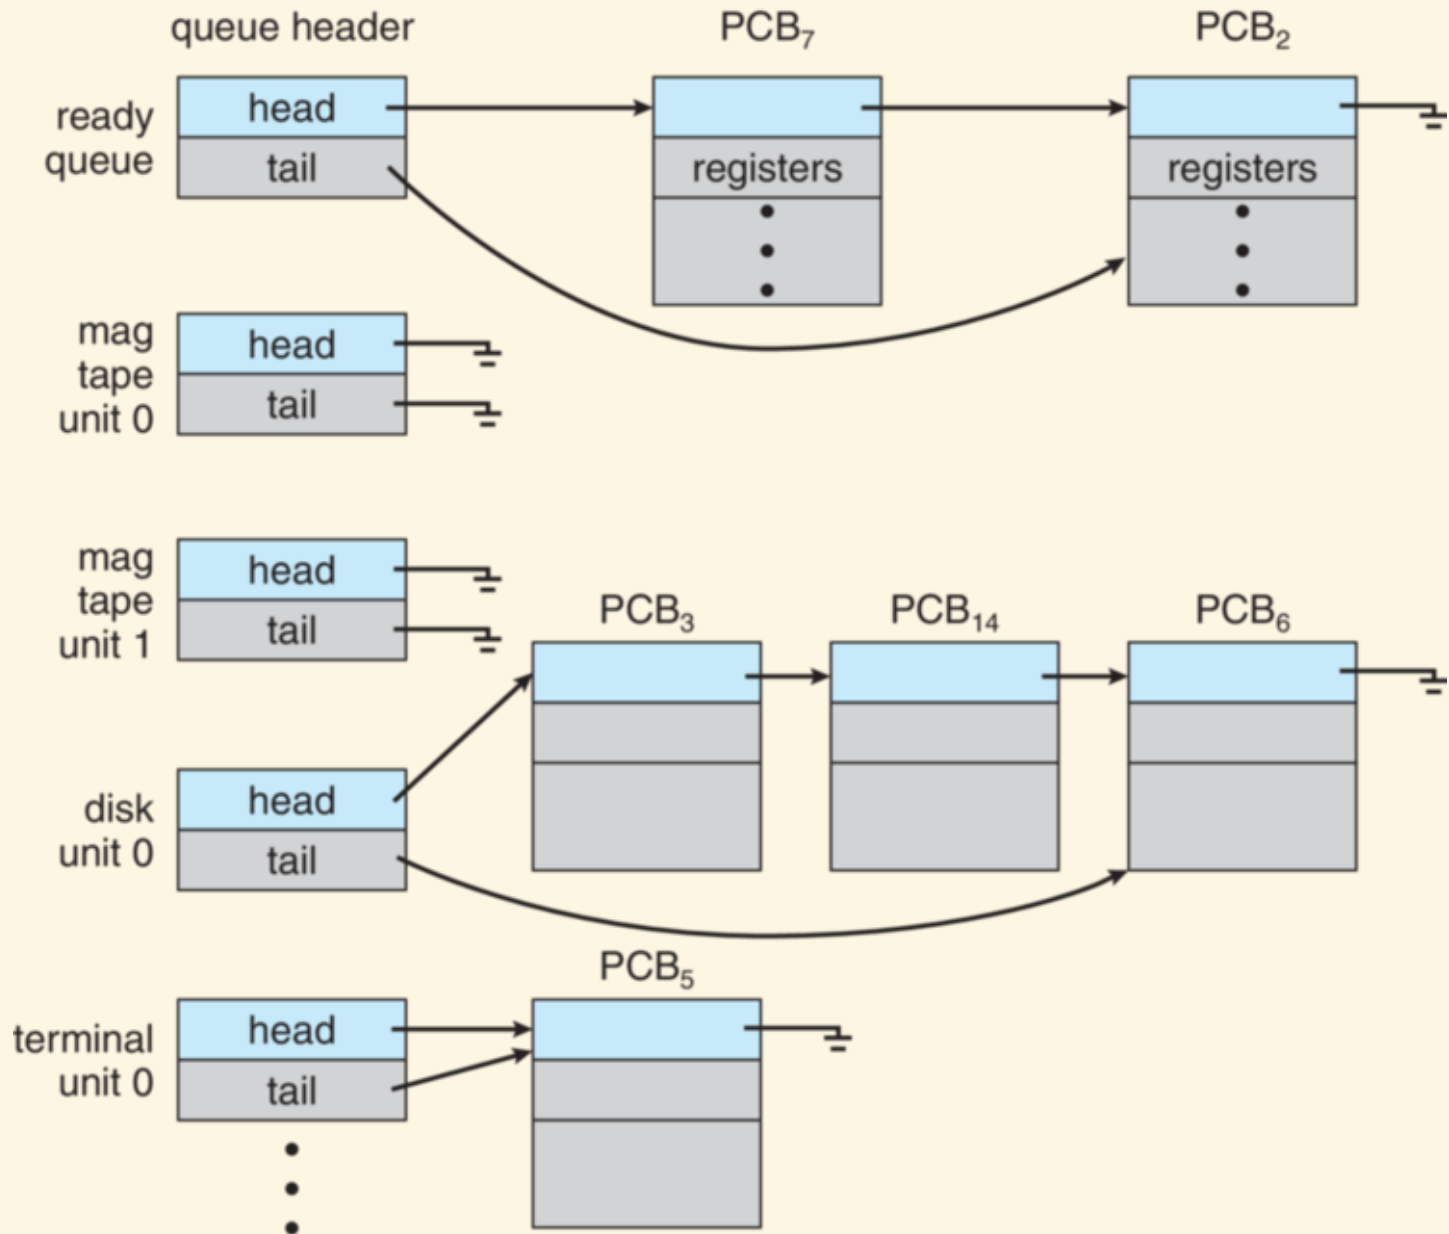
Nathan Bowman

Images taken from Silberschatz book

CPU Scheduling

We previously discussed the big picture of how processes are scheduled





We will now focus specifically on the CPU (short-term) scheduler

This decides which process from the ready queue will be run next

The ready queue is not necessarily a FIFO queue

Although we are discussing processes, the same ideas
apply to threads

We will discuss thread-specific issues later

We also focus on single-processor systems for now,
with details of multi-processor systems coming later

CPU scheduling is fundamental to achieving good resource utilization and system responsiveness

Want to ensure that there is some process using the CPU at all times

Scheduling objectives will depend on type of system, but some general rules apply

Key to effective scheduling: A given process will tend to alternate between **CPU bursts** and **I/O bursts**

CPU burst is when the program is actually performing useful work

Program will eventually need more data to operate on, so will make an I/O call and block

•
•
•

load store
add store
read from file

} CPU burst

wait for I/O

} I/O burst

store increment
index
write to file

} CPU burst

wait for I/O

} I/O burst

load store
add store
read from file


} CPU burst

wait for I/O

} I/O burst

•
•

CPU burst durations tend to be very short



New process may be assigned to CPU in any of four cases:

- running process switches from "running" to "waiting" (due to, e.g., an I/O call)
- running process switches from "running" to "ready" (due to, e.g., an interrupt)
- any process switches from "waiting" to "ready" (due to, e.g., completion of I/O)
- running process terminates

When running process blocks or terminates, new process should be chosen no matter what

In other two cases, different OSes may or may not even *consider* a context switch

Nonpreemptive scheduling (or cooperative scheduling)

-- OS will run scheduler only when running processes
blocks or terminates

Preemptive scheduling -- OS will run scheduler in any
of four previously mentioned cases

Modern general-purpose OSes tend to use preemptive
scheduling

Preemptive scheduling requires hardware support (a
timer)

Preemptive scheduling is partly to blame for the race conditions we observed previously

Some OSes will disable preemption during a system call to ensure the OS itself does not get into an inconsistent state

The **dispatcher** is responsible for actually giving control to scheduled process

Dispatcher must

- switch context
- switch to user mode
- jump to user program

Summary

- CPU scheduler determines order of processes in ready queue
- Effective scheduler must account for CPU-I/O burst cycle
- Scheduling may be preemptive or nonpreemptive