

# CIS 452 - Operating Systems Concepts


Nathan Bowman

Images taken from Silberschatz book

---

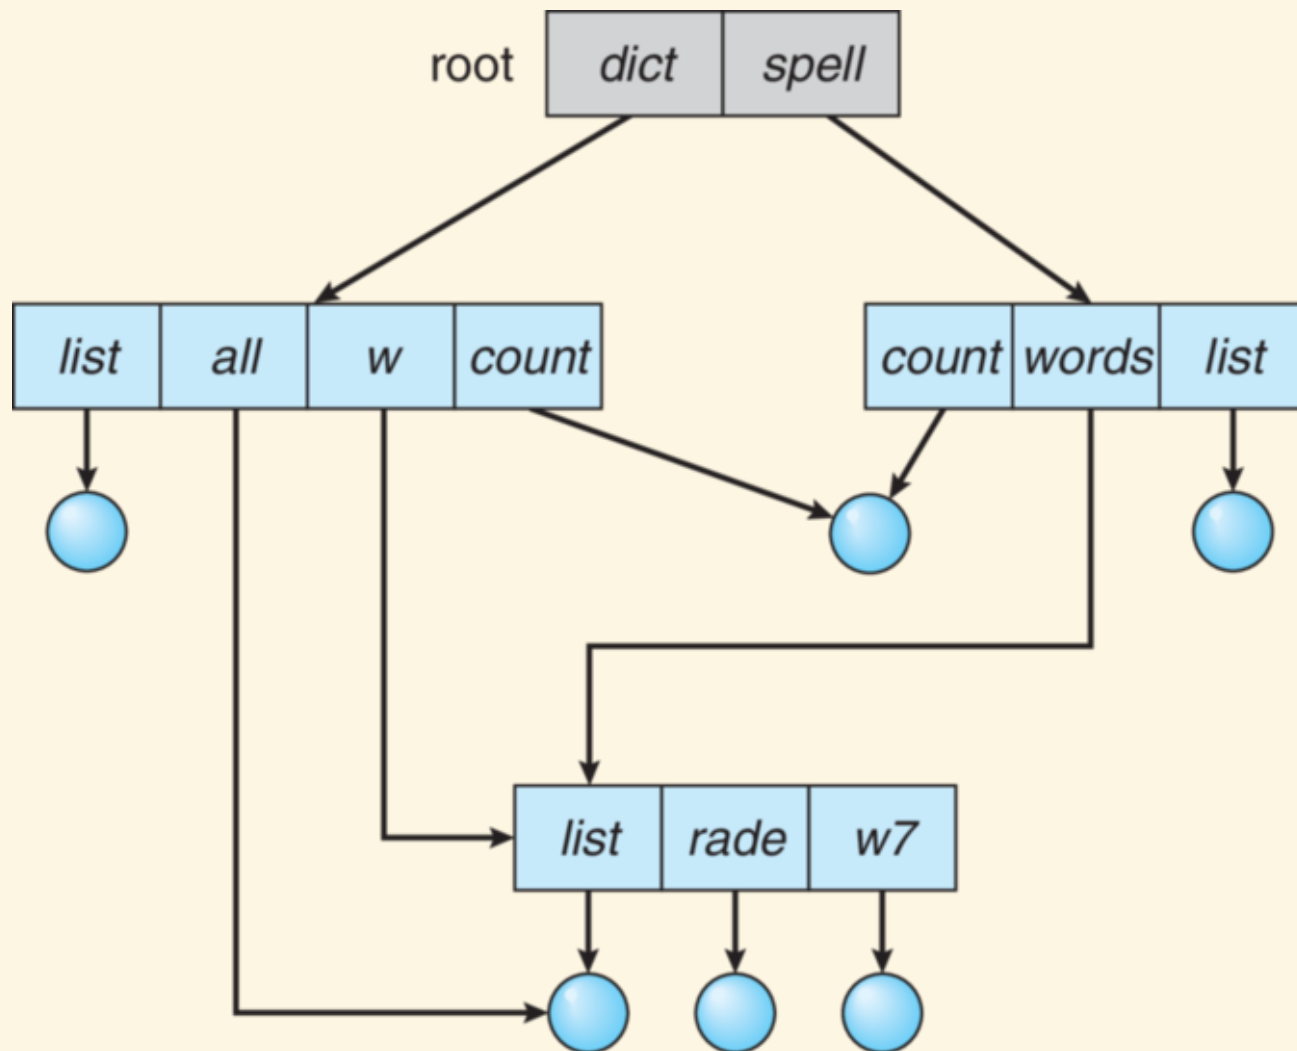
Links

Link exists when file or directory has more than one  
name



Links can be nonsymbolic (hard) or symbolic (soft)

**Nonsymbolic (hard) links** associate directory entry  
with another file



/dict

Name	inode
---	---
list	00000003
all	00000004
w	00000005
count	00000006

/spell

Name	inode
---	---
count	00000006
words	00000005
list	00000008

If `/dict/count` is updated, changes available immediately in `/spell/count`

Hard link is pointer to same underlying file (really, inode)

Note that names do not matter -- `/dict/w` is linked to `/spell/words`

Also, two or more files in same directory could be hard linked

/spell

Name	inode
---	---
count	00000006
words	00000005
list	00000008
list2	00000008



If files can be hard linked, OS must keep track of  
number of links to file

Space holding file cannot be reclaimed until all links  
removed

**Symbolic (soft) links** associate directory entry with another *path*

If `/spell/list2` were soft-linked to `/spell/list`, directory would look like

```
/spell
```

	Name		inode	
	---		---	
	count		00000006	
	words		00000005	
	list		00000008	
	list2		00000009	

<-- special link file

`list2` is a different file, but a special kind of file containing only filename `/spell/list1`

If program edits `list2`, it actually modifies whatever  
file resides at `/spell/list1`

What if different file is moved to `/spell/list1`?

```
/spell
```

	Name		inode	
	---		---	
	count		00000006	
	words		00000005	
	list		00000010	
	list2		00000009	

`list2` considers only the path

Changes to `list2` now affect inode 10 instead of inode  
8

Similarly, if `/spell/list1` is deleted, `list2` will be  
meaningless

Some programs do not "follow" symbolic links, meaning they act on link itself (e.g., act on inode 000000009 directly in below example)

See `man 7 symLink` for details of how this works in Linux

```
/spell
```

	Name		inode	
	---		---	
	count		00000006	
	words		00000005	
	list		00000008	
	list2		00000009	

# Hard link

---

```
$ ls -li  
12063523 -rw-r--r-- 1 nate nate 12 Nov 30 17:32 prog1.c
```

```
$ ln prog1.c hard.c
```

```
$ ls -li  
12063523 -rw-r--r-- 2 nate nate 12 Nov 30 17:32 hard.c  
12063523 -rw-r--r-- 2 nate nate 12 Nov 30 17:32 prog1.c
```

```
$ rm prog1.c
```

```
$ ls -li  
12063523 -rw-r--r-- 1 nate nate 12 Nov 30 17:32 hard.c
```

# Soft link

---

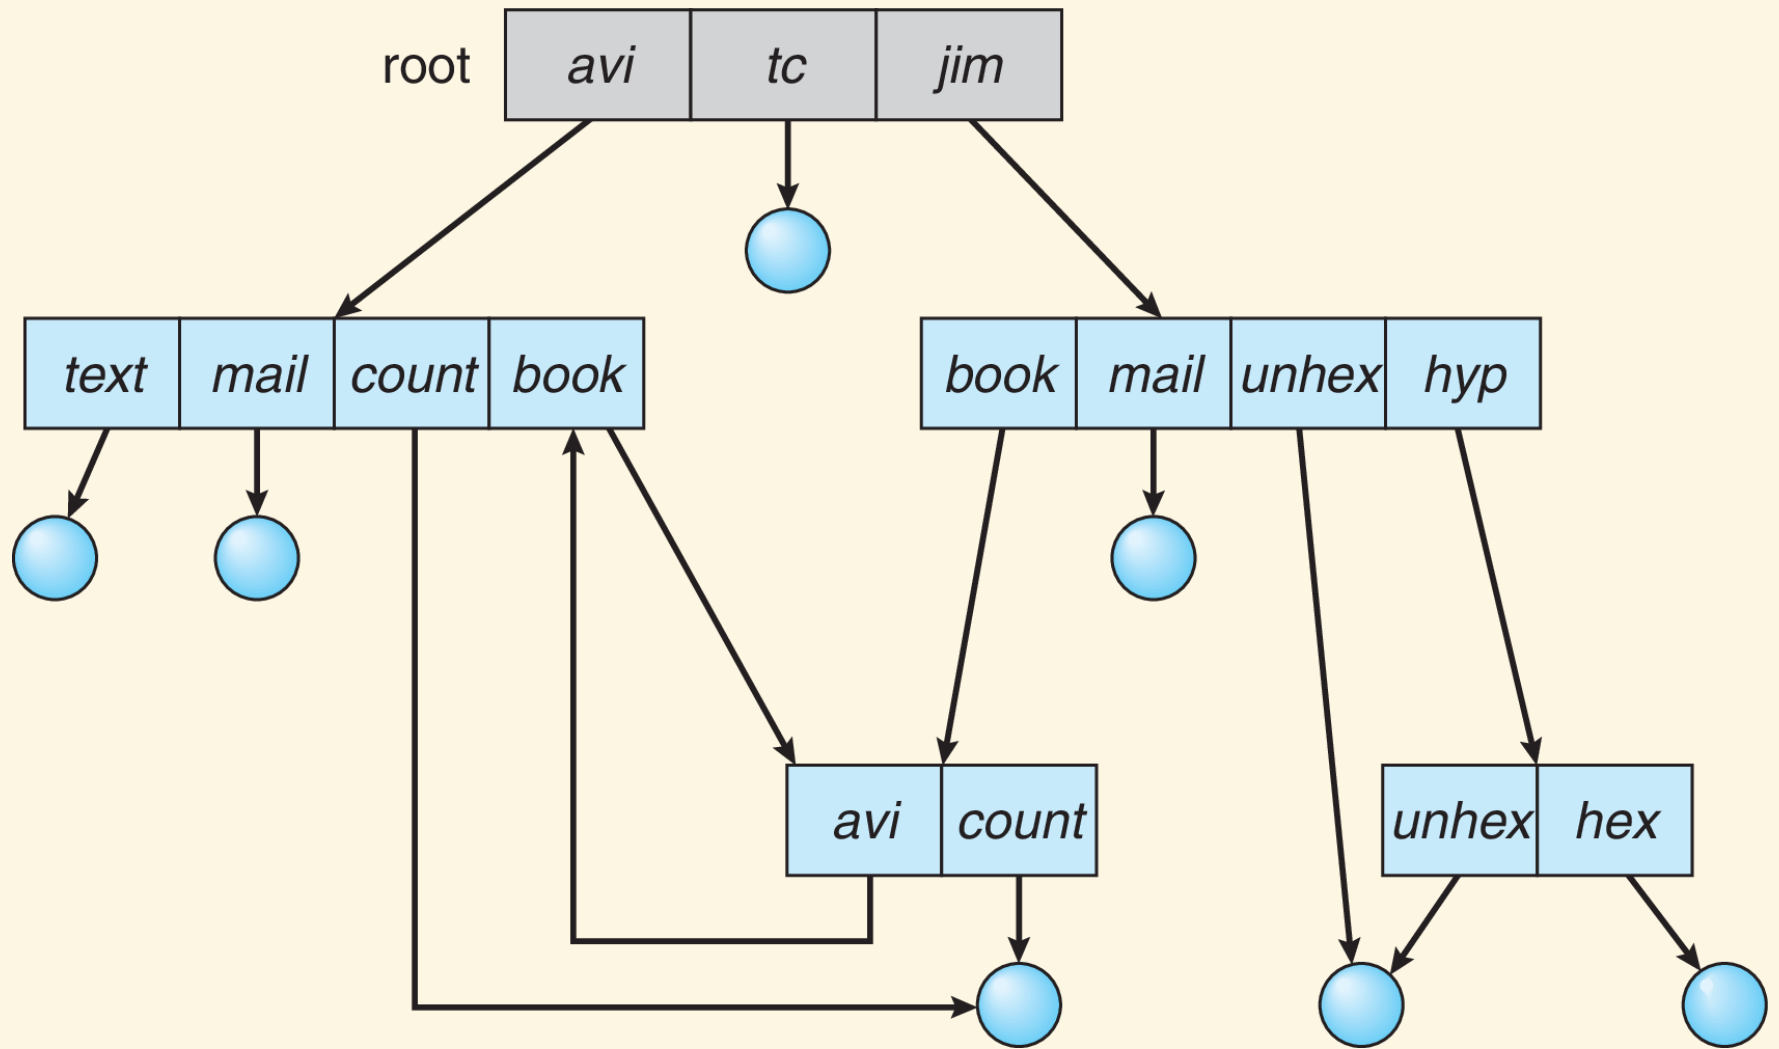
```
$ ls -li
12063523 -rw-r--r-- 1 nate nate 12 Nov 30 17:32 prog1.c

$ ln -s prog1.c soft.c
$ ls -li
12063523 -rw-r--r-- 1 nate nate 12 Nov 30 17:32 prog1.c
12063524 lrwxrwxrwx 1 nate nate 7 Nov 30 17:37 soft.c -> prog1.c

$ rm prog1.c
$ ls -li
12063524 lrwxrwxrwx 1 nate nate 7 Nov 30 17:37 soft.c -> prog1.c
```



Links are very useful, but they complicate structure



Problem with links, hard or soft, is that they can create cycles

Carelessly designed program can find itself in an infinite loop when traversing file tree

OS could disallow loops, but it would be expensive to run loop detection algorithm every time link is created

Linux disallows hard links to directories for this reason,  
but allows soft links to directories

Programs must be careful how they treat symbolic links  
when traversing file tree

See `man 7 symlink` for details